

Grouping Duplicate Citizen Complaints in a State Ombudsman

An operational account from Ceará, Brazil

FRANCISCO NAUBER BERNARDO GOIS, Controladoria e Ouvidoria Geral do Estado do Ceará (CGE-CE), Brazil

BERG SILVA, Controladoria e Ouvidoria Geral do Estado do Ceará (CGE-CE), Brazil

A citizen can tell the same story ten times in one morning. Each wording looks like a new file. The officer reads the pile twice, and the person who wrote waits.

Similaridade is the grouping job used by the Controladoria e Ouvidoria Geral do Estado do Ceará (CGE-CE). It asks whether two free-text complaints point to the same event. On one historical batch of 4,389 real complaints, the production stack (TF-IDF + DBSCAN, keep a group if mean intra-cluster cosine similarity ≥ 0.75) grouped 62.4% of the texts. The same batch, encoded with sentence embeddings and clustered with HDBSCAN, grouped an inferred 82.5% before any language model. Adding a leftover pass through DeepSeek R1 raised the reported share to 84.9% (3,725 complaints). That last step recovered 104 of 768 leftovers after seeing only 27 candidates. How those 27 were chosen is not in the team note. Whether the model ran inside CGE is also unknown. We therefore treat 84.9% as an ablation with an undocumented filter, not as a deployment result.

None of these rates is a precision. The Ouvidoria has not labeled a sample of “same case.” In September 2026 the live job at times collapsed to a single cluster. The claim that survives is smaller: lexical clustering already catches copy-paste waves; a semantic stack grouped more of the same historical set; daily production of that stack, and any leftover language-model path, is still being earned.

CCS Concepts: • **Applied computing** → **Computing in government**; • **Information systems** → *Clustering and classification*; • **Computing methodologies** → *Natural language processing*.

Additional Key Words and Phrases: citizen complaints, ombudsman, near-duplicate detection, text clustering, sentence embeddings, HDBSCAN, digital government, Brazil

Draft v0.5 — 24 September 2026. Practice paper aimed at ACM DGOV, not submitted. Authorship is provisional. Diagrams: Grok Imagine, type set from team counts. Prior Sonar read: docs/perplexity-review.md. Venue note: docs/journal.md.

1 INTRODUCTION

An ombudsman file starts as free text. The same broken pipe, the same delay at a public counter, or the same coordinated campaign can arrive tens of times before lunch. If each arrival becomes its own case, the officer repeats the same reading and the citizen waits longer (Figure 1).

Similaridade sits on the complaint stream of the state Ouvidoria (COUVI). The job is narrow. Put texts that likely refer to the same event in one group. Mark the rest as leftovers. Show the officer a provisional review queue—green, yellow, red—instead of a raw list.

Most public-sector NLP work classifies a complaint into a department or extracts a theme [5, 13, 18, 20]. Similaridade is closer to near-duplicate detection [2, 12]: it asks whether two texts are the same case. TF-IDF, DBSCAN, sentence transformers, and HDBSCAN are known tools [3, 8, 14, 16]. What this draft can still offer is an operational account: one Brazilian state stream, three treatments on the same 4,389 complaints, a conservation table that does not fully close, and a protocol for the labels that are still missing.

Four questions stay after the August tests.

(1) How much of a real complaint batch does lexical clustering already group?

Authors' addresses: Francisco Nauber Bernardo Gois, Controladoria e Ouvidoria Geral do Estado do Ceará (CGE-CE), Fortaleza, Brazil, francisco.gois@cge.ce.gov.br; Berg Silva, Controladoria e Ouvidoria Geral do Estado do Ceará (CGE-CE), Fortaleza, Brazil.




Fig. 1. Graphical abstract. Scene drawn with Grok Imagine; title, rates, and the limit set in type from the team note. Grouping rate is not precision. The leftover language-model step remains an ablation.

- (2) How much more does embedding + HDBSCAN group on the same batch, before any language model?
- (3) What does the leftover language-model pass add, and what must be known before that add is credited?
- (4) What still breaks when the job runs every day?

This version does not invent officer labels, clustering radii, or a model name. Where the team note is silent, the table says unknown.

The first venue for this draft is ACM *Digital Government: Research and Practice* (DGOV), practice track. The paper is an operational account, not a new clustering theorem. Method journals would ask for a gold partition we do not have.

2 STATE OF THE ART

Public grievance research splits into three jobs that are easy to mix up. Figure 2 is the cut that matters for an ombudsman officer.

Routing and classification. A first line assigns a complaint to a unit or a service. Das et al. used sentence-level sentiment to help officers triage incoming text [5]. Rakhimzhanov, Belginova, and Yedilkhan compared large




Fig. 2. Three jobs that public-sector NLP often mixes. Route sends a text to a unit. Theme sorts a mixed pile. Case puts two tellings of one event in one folder. Only the third job is Similaridade. Diagram: Grok Imagine; labels set in type.

language models with instruction-tuned embeddings on municipal transport complaints; fine-tuned embeddings reached about 90% accuracy and were cheaper than prompting an LLM for every item [13]. Silva et al. linked Brazilian federal ombudsman manifestations on Fala.BR to official public services with BM25 and dense retrieval plus RAG-expanded labels [18]. That is routing. Two texts about the same water outage can share a service and still be two cases, or two services and still be one case.

Theme discovery. A second line asks what citizens talk about. BERTopic embeds documents, reduces them, clusters with HDBSCAN, and labels each cluster with class-based TF-IDF [10]. Tang, Pan, and Gu used that recipe on China’s online government inquiry platform [20]. Theme models help policy. They are the wrong metric when an officer must link a copy-paste wave that shares one event and many wordings.

Duplicate and same-event grouping. A third line asks whether two records are the same complaint or the same event. That is record linkage and duplicate detection, not topic modelling [6, 9]. Broder’s resemblance test did the job for web pages [2]. Rabbi et al. did it for Portuguese consumer complaints across Consumidor.gov.br, Procon, and Sindec: BERTimbau embeddings [19], a cosine threshold set with domain experts (90% in their study), and a 30-day window. Most of their duplicates arrived inside that window [12]. Their unit is often the same consumer repeating a claim. An ombudsman wave can also be many people writing one outage. Esperança et al. clustered informatics complaints in a Portuguese public agency so repeated issues surface once; the target is still a category [7]. Romberg and Escher, writing in DGOV itself, treat duplicate detection as one of three evaluation jobs on citizen text, and they name the same gaps we still have: quality of the groups, non-English corpora, and a tool an officer can actually run [15]. Chen et al. turn government complaints into a classification problem so a server can route a reply [4]. That is still routing. Similaridade does not pick a department. It asks whether two files should share one case.

Table 1. Public-sector complaint work versus Similaridade. Theme = latent topic. Route = department or service. Case = same event.

Work	Task	Representation	Grouping
Fellegi and Sunter [9]	Record link	Field agreement	Likelihood ratio
Elmagarmid et al. [6]	Duplicate record	Survey	Pair classifiers
Romberg and Escher [15]	Dupl. / theme / argument	Review (DGOV)	Mixed NLP
Chen et al. [4]	Route / reply	Text classifier	Supervised
Das et al. [5]	Route / triage	Sentiment features	Supervised
Tang et al. [20]	Theme	Embeddings	BERTopic
Rakhimzhanov et al. [13]	Route	E5 / BGE / LLM	Classifier
Silva et al. [18]	Route (service)	BM25 + dense + RAG	Retrieval
Esperança et al. [7]	Theme / pattern	BERT embeddings	Semantic clusters
Rabbi et al. [12]	Duplicate	BERTimbau	Threshold + 30 days
Similaridade (production)	Case	TF-IDF	DBSCAN
Similaridade (no LLM)	Case	Sentence embeddings	HDBSCAN
Similaridade (leftover LLM)	Case / theme?	Prompt on leftovers	Filter unknown

Where *Similaridade* sits. Table 1 places the stacks against that map. Production Ceará is lexical near-duplicate detection on short state-ombudsman texts, without the identity keys and time window Rabbi et al. could use. The homolog stack is embeddings + HDBSCAN, plus an optional leftover language model. It is not a BERTopic run: there is no class-based TF-IDF topic label on the main clusters. What this draft still lacks is a human label of “same case.” Without that label, a higher grouping rate can just mean a more permissive merge. The metric that would decide the argument is BCubed precision and recall against officer gold [1], not the share of texts placed in some group.

On the method side the genealogy is short. Term weights still catch identical paste [16]. DBSCAN finds dense regions and treats sparse points as noise [8, 17]. That matches an ombudsman day: one campaign, or many singletons. The price is a fixed radius. HDBSCAN drops that radius [3, 11]. Sentence-BERT gives vectors that cosine can compare without a cross-encoder on every pair [14]. Portuguese checkpoints exist [19]. This draft does not yet name the checkpoint that ran.

3 SETTING

CGE-CE is the state controller and ombudsman of Ceará. Similaridade serves the Ouvidoria. ASES builds and runs the software. Airflow starts the job. Counts later land in an internal log table (etl.similaridade_log: complaints processed, grouped, leftovers, largest and smallest cluster, intra-cluster similarity).

The production schedule in August 2026 was three runs per day: 08:00, 12:00, and 17:00. Source code lives on the state’s GitLab, project `analise_similaridade_dev`. This public draft does not copy that repository.

The 4,389-complaint homolog batch was reported to the team on 28 August 2026. The note does not give the extraction window, language mix, length histogram, or agency mix. Complaint text is sensitive. We report only aggregate counts already shared inside the team. No complaint body is reproduced.

4 DEFINITIONS AND COUNTING

The words below are operational, not legal.

- **Complaint.** One incoming free-text record in the batch.

Table 2. Conservation of the 4,389-complaint homolog batch. Production is reported only as a rate; 2,739 is $0.624 \times 4,389$, rounded. The last two homolog rows are implied, not independently logged.

Quantity	Production	Homolog
Incoming complaints	4,389	4,389
Grouped (reported)	$\approx 2,739$ (62.4%)	3,725 (84.9%)
Accepted clusters (reported)	—	354
Leftovers before LLM (reported)	—	768
Candidates sent to LLM (reported)	—	27
Recovered leftovers (reported)	—	104
HDBSCAN-only grouped (implied)	—	3,621 (82.5%)
Leftovers after LLM (implied)	—	664

- **Cluster.** A set of complaints that the algorithm labeled together.
- **Grouped.** A complaint in an accepted cluster of size ≥ 2 .
- **Leftover.** A complaint not grouped after that stack’s clustering step.
- **Candidate.** A leftover passed to the language model.
- **Recovered.** A leftover that the language model placed in a group the team later counted as grouped.
- **Theme (leftover path).** The team’s word for a leftover group returned by the language model. It may be a topic, not a same-event case.
- **Band.** A provisional review queue (green / yellow / red). Not a measured safety threshold.

Table 2 is the conservation we can write. The reported homolog grouped count is 3,725 of 4,389 (84.9%). That number includes 104 recovered leftovers. Remove them and HDBSCAN alone grouped 3,621 (82.5%). Two facts still do not close.

- (1) **Where the 21 leftover groups sit.** Scenario A: they are inside the 354 clusters ($354 - 21 = 333$ HDBSCAN clusters). Scenario B: they sit beside the 354 clusters ($354 + 21 = 375$ groups after the language model). The team note does not choose.
- (2) **The 27-to-104 step.** 27 candidates cannot contain 104 recovered complaints unless each candidate stands for a bundle, or the 27 are seeds that pull neighbors, or the note compressed a longer list. None of those rules is written down. Until the rule is pinned, the leftover pass is not an auditable ablation.

The 17:00 production run on 28 August 2026 is not in the team note. The two reported runs that day add to the stated day total ($49 + 78 = 127$).

5 PROBLEM STATEMENT AND FORMULATION

Let $X = \{x_1, \dots, x_n\}$ be a batch of free-text complaints. An unknown map $e : X \rightarrow \mathcal{E}$ sends each text to the real-world event the citizen is talking about. Two texts are the same case when they share that event:

$$x_i \sim x_j \iff e(x_i) = e(x_j). \quad (1)$$

The relation $\sim$ is an equivalence. Its classes are the gold files an officer would open. We do not observe e . Similaridade builds a predicted partition $\widehat{\Pi}$ of X that tries to recover those classes, then leaves the stamp to a human.

This is record linkage on one file: deduplication of event mentions, not of people [6, 9]. Fellegi and Sunter decide, for each pair, link / non-link / possible link, at stipulated error rates. We first cluster, then apply a similar three-way queue to each predicted class (Section 5.4). Pairwise classification of all $\binom{n}{2}$ pairs is the exact problem; density clustering is the scalable surrogate.

5.1 Representation

An encoder φ sends each text to the unit sphere,

$$\varphi : \mathcal{X} \rightarrow \{v \in \mathbb{R}^d : \|v\|_2 = 1\}. \quad (2)$$

Production uses TF-IDF on unigrams to trigrams [16]. Homolog uses a sentence embedding [14]. Cosine similarity and cosine distance are

$$s_{ij} = \langle \varphi(x_i), \varphi(x_j) \rangle, \quad d_{ij} = 1 - s_{ij}. \quad (3)$$

On the sphere, $s_{ij} \in [-1, 1]$. Short texts make $\varphi(x_i)$ a poor object: the vector is defined, the neighborhood is not.

5.2 Density clustering

DBSCAN treats i as a core point when its ε -neighborhood is large enough [8, 17]:

$$N_\varepsilon(i) = \{j : d_{ij} \leq \varepsilon\}, \quad i \text{ is core} \iff |N_\varepsilon(i)| \geq m. \quad (4)$$

A predicted cluster C is a maximal density-connected set of core points and their border points. Points that are not density-reachable are leftovers L . The production note does not record ε or m . HDBSCAN replaces the single radius ε by a hierarchy of density levels and cuts that tree [3, 11]. The homolog note does not record the cut.

5.3 Acceptance and the grouping rate

A cluster is kept only if it is not a singleton and its mean pairwise cosine clears a threshold $\tau = 0.75$:

$$\bar{s}(C) = \frac{2}{|C|(|C| - 1)} \sum_{\substack{i, j \in C \\ i < j}} s_{ij}, \quad C \text{ is accepted} \iff |C| \geq 2 \wedge \bar{s}(C) \geq \tau. \quad (5)$$

Let $\mathcal{A}$ be the family of accepted clusters. The grouped set and the grouping rate are

$$G = \bigcup_{C \in \mathcal{A}} C, \quad \rho = \frac{|G|}{n}. \quad (6)$$

The numbers in this draft are values of ρ , not of precision. A more permissive (ε, m, τ) raises ρ by merging unrelated events. That is why $\rho_{\text{prod}} = 0.624$ and $\rho_{\text{HDBSCAN}} = 0.825$ on the same $n = 4389$ cannot be ranked without gold.

Mean cosine is not transitive. There exist accepted clusters in the production log where s_{AB} and s_{BC} are high and $s_{AD} \approx 0.006$. The chain sits inside one label; the pair (A, D) would not pass τ by itself. A later check flags an intruder by a z-score on the distance to the centroid $\mu_C = \frac{1}{|C|} \sum_{i \in C} \varphi(x_i)$:

$$\delta_i = \|\varphi(x_i) - \mu_C\|_2, \quad z_i = \frac{\delta_i - \text{mean}(\delta_C)}{\text{sd}(\delta_C)}. \quad (7)$$

The team reported that this check behaved better on embeddings than on TF-IDF. The cut on z_i is not in the note.

5.4 Leftovers and the three queues

Write $L = \mathcal{X} \setminus G$ after clustering. The homolog leftover pass selects a candidate set $K \subset L$ with $|K| = 27$ out of $|L| = 768$ and returns a recovered set R with $|R| = 104$. The selection map $L \mapsto K$ is unknown, so the step is not an auditable operator. Until it is written down we report

$$\rho_{\text{HDBSCAN}} = \frac{|G \setminus R|}{n} = \frac{3621}{4389}, \quad \rho_{\text{HDBSCAN}+R} = \frac{|G|}{n} = \frac{3725}{4389} \quad (8)$$

as two different treatments, not as one system.




Fig. 3. Homolog path on the same incoming text: complaint, embed, HDBSCAN, leftover, officer stamp. The leftover station is an ablation; host and candidate rule are unknown. Diagram: Grok Imagine; labels set in type.

Accepted clusters are shown in three queues that replay Fellegi and Sunter's three decisions [9]: green $\approx$ link, yellow $\approx$ possible link, red $\approx$ read as if non-link. The cuts that send a cluster to a queue are unknown. About 87% of $|G|$ sat in green. The arithmetic remainder $0.13 \times 3725 \approx 484$ is not a measured officer load.

5.5 What a later evaluation must compute

Let Π^* be the gold partition induced by e . Against a blinded officer sample of Π^* , the scores that decide the argument are BCubed precision and recall [1], the false-merge rate (two events in one $\widehat{C}$), and the false-split rate (one event in two $\widehat{C}$). Pairwise precision of accepted pairs inside green is the safety number for the bulk-link button. ρ stays in the log. It is not the headline metric of a submission.

6 METHODS

Figure 3 shows the two stacks. Table 3 is the experiment manifest. Empty cells are not omissions of convenience. They are missing from the source note.

6.1 Production stack (TF-IDF + DBSCAN)

Each complaint is a document. The encoder is TF-IDF over unigrams, bigrams, and trigrams [16]. Vectors are L_2 -normalized. DBSCAN then groups points that share a neighborhood of a given radius [8]. After clustering, a group is accepted only if mean pairwise cosine similarity inside the group is at least 0.75.

Similarity is not transitive. A can sit close to B , B to C , and C to D , while A and D score near 0.006 and still share a label. A later check uses distance to the cluster centroid, as a z-score, to flag intruders. That check behaved better on sentence embeddings than on TF-IDF.

6.2 Homolog stack (embeddings + HDBSCAN)

Encode each complaint with a Sentence-Transformers model [14]. Cluster with HDBSCAN [3]. The checkpoint name is not in this file. Portuguese models such as BERTimbau exist [19]; naming one of them would be a guess.

Table 3. Experiment manifest. “Unknown” means the team note does not record the value. A later frozen run must fill every unknown cell before the homolog stack is treated as reproducible.

Item	Production	Homolog
Batch size	4,389 (same set)	4,389
Batch window / checksum	Unknown	Unknown
Encoder	TF-IDF, 1–3 grams [16]	Sentence-Transformers, name unknown
Vector norm	L_2	Unknown
Clusterer	DBSCAN [8]	HDBSCAN [3, 11]
Radius / min. cluster size	Unknown	Unknown
Acceptance rule	Mean pairwise cosine ≥ 0.75	Unknown if the same cut applies
Language model	None	DeepSeek R1; date, prompt, host unknown
Candidate rule	—	Unknown (768 $\rightarrow$ 27)
Software pin	Unknown	Worker later loaded ST 5.1.2 vs. model 5.4.0
Hardware / seed	Unknown	Embed step $\approx$ 3 min on $\approx$ 4,000 texts
Officer gold	None	None

Table 4. Grouping rates on the 4,389-complaint batch. The last column is not precision.

Treatment	What the note supports	Grouped
TF-IDF + DBSCAN	Production stack, live in August 2026	62.4%
Embeddings + HDBSCAN	Implied by removing 104 recoveries	82.5%
HDBSCAN + leftover LLM	Reported homolog total; filter unknown	84.9%

6.3 Leftover language-model pass (ablation, not a deployment claim)

The team sent a reduced leftover list to DeepSeek R1 and counted 104 recoveries in 21 groups. Because the candidate rule and the hosting path are unknown, this pass is reported as an ablation and is withdrawn from the operational claim. A later run that sends leftover text off-site needs a written Ouvidoria basis, a retention limit, and a log of what left the building.

6.4 Provisional review queues

Grouped items are shown in three queues (Figure 4): green (suggested bulk linking), yellow (officer confirms), and red (officer reads with care). A cluster at 100% internal similarity means the texts are identical. About 87% of grouped homolog items sat in green. If that queue were trusted, the remaining 13% of 3,725 would be about 484 complaints. That 484 is an arithmetic remainder. It is not officer minutes, and the band cut itself is not in the note.

7 RESULTS

7.1 Three treatments, one batch

Figure 5 and Table 4 split what v0.3 still printed as one jump.

The cleaner comparison is 62.4% versus 82.5%. That is still only a grouping-rate gap of 20.1 percentage points on one historical batch. A more permissive merge would also raise the rate. The leftover language model adds 2.4 points only if the 104 recoveries are accepted at face value.

The leftover path is in Figure 6: 768 $\rightarrow$ 27 $\rightarrow$ 104 in 21 groups the team called themes.



Fig. 4. Conceptual illustration (Grok Imagine). Three trays: suggested bulk link, confirm, read with care. The stamp waits on every tray. These are interface queues, not a measured safety threshold.

7.2 One production day

The live TF-IDF job on 28 August 2026 shows what lexical clustering does when a wave arrives (Figure 7). At 08:00, 48 of 49 complaints sat in one cluster (mean similarity 85.68%). At 12:00, 73 of 78 sat in a cluster of identical texts. The day total was 127 processed, 124 grouped (97.6%), 3 leftovers.

A 97.6% grouping rate on a campaign day is not in conflict with 62.4% on the long set. The long set mixes ordinary days and waves. 28 August was a wave day.

7.3 What broke after the test

September 2026 showed the gap between a homolog notebook and a daily Airflow job. A homolog build of the new stack failed and later recovered. The live grouping collapsed toward one cluster and many leftovers. Infrastructure blocked Python library downloads. The Airflow worker loaded Sentence-Transformers 5.1.2 against a model saved with 5.4.0. Short texts carried little signal, and cosine grouping then failed. So 82.5% and 84.9% are batch results. They are not the steady daily rate.

8 HUMAN VALIDATION PROTOCOL (UNRUN)

No COUVI officer has labeled this batch. The protocol below is what v0.5 must run. It is not a result.

Unit. A pair or a cluster is *same event* when a reasonable officer would open one case file. *Related theme* means the same subject without the same event. *Different* and *unclear* are the other two codes.



Fig. 5. Three treatments on the same 4,389 complaints. Stack height is schematic, not to scale. Printed rates are the team counts. The middle value is inferred (3,725 – 104). The right-hand 84.9% includes an undocumented leftover filter. None of the rates is a precision. Diagram: Grok Imagine; numbers set in type from the team note.

Sample. Stratify. About 50 green clusters, 30 yellow, 20 red, 40 leftovers, and 30 pairs drawn from clusters larger than 20. Include short texts (under 20 tokens) as their own stratum if the log can list them. Two officers, blinded to stack and to each other. A third officer adjudicates disagreements.

Metrics. Pairwise precision of proposed merges; false-merge and false-split rates; Cohen’s κ ; BCubed precision and recall against the adjudicated clustering [1]. Report each treatment separately: production, HDBSCAN, HDBSCAN plus leftover model—and only after the leftover filter is written down.

Workload. If the queues go into a live week, log accept, reject, split, and minutes per 100 complaints. Until those minutes exist, do not call 484 an officer load.

9 DISCUSSION

Lexical clustering already catches the easy wave: identical or near-identical paste. That is useful on a morning like 28 August. It is the wrong tool when two people write the same event in different words.

The homolog HDBSCAN step moves the cut from tokens to sense. Rabbi et al. needed a time window and identity keys to call two consumer texts a duplicate [12]. Similaridade does not yet use those keys. That is a gap, not a virtue: a 30-day window and an agency field would be the first cheap control on short texts.

The leftover language-model pass follows the cost lesson in Rakhimzhanov et al. [13]—spend the expensive model on leftovers—but the 768 → 27 → 104 funnel is not yet a result. If the 27 candidates were chosen by a person who already saw the 104, the credit is circular.

Compared with Silva et al. [18] and Tang et al. [20], the claim here is not a new router and not a dashboard of themes. It is a count of how many ombudsman files the machine put in a group before a human reads them.




Fig. 6. Conceptual illustration (Grok Imagine) of the leftover path. Counts on this batch: $768 \rightarrow 27 \rightarrow 104$ in 21 groups. The selection rule for the 27 is unknown, so the 104 cannot yet be audited.

Fellegi and Sunter gave the pair decision [9]. Broder gave resemblance for pages [2]. The ombudsman pile is shorter, in Portuguese, and closed by a stamp.

10 LIMITATIONS

The limits are part of the result.

- (1) Counts come from team reports, not from a frozen public extract of `etl.similaridade_log`.
- (2) There is no officer gold, so there is no precision, recall, κ , or BCubed score.
- (3) The embedding checkpoint, DBSCAN radius, HDBSCAN size, and band cuts are unknown.
- (4) The leftover language-model pass has an undocumented filter and an unstated host.
- (5) September failures show that the homolog rates are not the current daily truth.
- (6) We do not measure minutes saved for the officer or days saved for the citizen.
- (7) Authorship is incomplete.

A submission without the protocol in Section 7, filled, will still read as a slide.

11 ETHICS AND DATA

Complaint text can identify a person or a public servant. This draft uses only counts. Any later replication package must stay inside CGE, with access logged, or use a redacted sample approved by the Ouvidoria.

The leftover pass used DeepSeek R1. The team note does not say whether that call stayed on a machine inside CGE. Until that path is written down, we do not treat the leftover pass as part of the production design. A later




Fig. 7. Production TF-IDF + DBSCAN on 28 August 2026. Large clusters read as copy-paste or light paraphrase, not as ordinary mixed traffic.

run that sends leftover text off-site needs a written Ouvidoria basis, a retention limit, a ban on provider training use, and a log of what left the building.

Clustering is an aid. An officer still closes the case. Linking two files that are not the same event can bury a second victim. That is why green stays a queue.

12 CONCLUSION

On 4,389 real complaints, TF-IDF + DBSCAN grouped 62.4% of the texts. Embeddings + HDBSCAN grouped an inferred 82.5% before any language model. A leftover pass then raised the reported share to 84.9%, but the filter and the host of that pass are unknown, so that last step is not a deployment claim. A wave morning in August showed the lexical stack at work: 97.6% grouped that day, almost all in two large clusters.

The next version needs what this draft cannot invent: COUVI labels on a stratified sample, a filled manifest, and a week of production logs after Airflow is stable. Until then the honest claim is smaller. The Ouvidoria already groups copy-paste waves. A semantic stack grouped more of the same historical batch. Daily production of that stack is still being earned.

ACKNOWLEDGMENTS

ASESI colleagues who discussed the method in June–September 2026, the COUVI officers who will have to trust or reject the green queue, and Lucas Pimentel, who hosts the internal GitLab project. Errors in this draft are the authors’.

CREDIT (PROVISIONAL)

Conceptualization and writing (original draft): Francisco Nauber Bernardo Gois. Methodology, software, and investigation (homolog measurements): Berg Silva. Other roles to be assigned.

FUNDING

No external grant. Work done as part of CGE-CE / ASES operations.

COMPETING INTERESTS

The authors work for the agency that operates the system. That is the point of a practice paper. It is also a conflict to declare.

DATA AVAILABILITY

Complaint text is not public. Aggregate figures can be checked against CGE internal logs by staff with access.

USE OF AI TOOLS

A language model helped organize this English draft from team notes. Diagrams and the graphical abstract were drawn with Grok Imagine (model grok-imagine-image-quality). Titles, rates, and captions were set in type from the team counts; the generator was not asked to write numerals. The 28 August production-day chart remains matplotlib. No complaint text was sent to an image or language model. The corresponding author is responsible for every number.

REFERENCES

- [1] Enrique Amigó, Julio Gonzalo, Javier Artiles, and Felisa Verdejo. 2009. A comparison of extrinsic clustering evaluation metrics based on formal constraints. *Information Retrieval* 12, 4 (2009), 461–486. <https://doi.org/10.1007/s10791-008-9066-8>
- [2] Andrei Z. Broder. 1997. On the Resemblance and Containment of Documents. In *Proceedings of the Compression and Complexity of Sequences 1997*. IEEE, 21–29. <https://doi.org/10.1109/SEQUEN.1997.666900>
- [3] Ricardo J. G. B. Campello, Davoud Moulavi, and Jörg Sander. 2013. Density-Based Clustering Based on Hierarchical Density Estimates. In *Advances in Knowledge Discovery and Data Mining (LNCS, Vol. 7819)*. Springer, 160–172. https://doi.org/10.1007/978-3-642-37456-2_14
- [4] Siqi Chen, Yanling Zhang, Bin Song, Xiaojiang Du, and Mohsen Guizani. 2022. An Intelligent Government Complaint Prediction Approach. *Big Data Research* 30 (2022), 100336. <https://doi.org/10.1016/j.bdr.2022.100336>
- [5] Rajesh Kumar Das, Manas Panda, and Harekrishna Misra. 2020. Decision Support Grievance Redressal System Using Sentence Sentiment Analysis. In *Proceedings of the 13th International Conference on Theory and Practice of Electronic Governance*. ACM, 17–24. <https://doi.org/10.1145/3428502.3428505>
- [6] Ahmed K. Elmagarmid, Panagiotis G. Ipeirotis, and Vassilios S. Verykios. 2007. Duplicate Record Detection: A Survey. *IEEE Transactions on Knowledge and Data Engineering* 19, 1 (2007), 1–16. <https://doi.org/10.1109/TKDE.2007.9>
- [7] Marco Esperança, D. Cuoghi de Carvalho Veríssimo Freitas, Pedro V. Paixão, Tomás A. Marcos, Rafael A. Martins, and João C. Ferreira. 2025. Proactive Complaint Management in Public Sector Informatics Using AI: A Semantic Pattern Recognition Framework. *Applied Sciences* 15, 12 (2025), 6673. <https://doi.org/10.3390/app15126673>
- [8] Martin Ester, Hans-Peter Kriegel, Jörg Sander, and Xiaowei Xu. 1996. A Density-Based Algorithm for Discovering Clusters in Large Spatial Databases with Noise. In *Proceedings of the Second International Conference on Knowledge Discovery and Data Mining*. AAAI Press, 226–231. <https://cdn.aaai.org/KDD/1996/KDD96-037.pdf>
- [9] Ivan P. Fellegi and Alan B. Sunter. 1969. A Theory for Record Linkage. *J. Amer. Statist. Assoc.* 64, 328 (1969), 1183–1210. <https://doi.org/10.1080/01621459.1969.10501049>
- [10] Maarten Grootendorst. 2022. BERTopic: Neural Topic Modeling with a Class-Based TF-IDF Procedure. *arXiv preprint arXiv:2203.05794* (2022). <https://doi.org/10.48550/arXiv.2203.05794>
- [11] Leland McInnes, John Healy, and Steve Astels. 2017. hdbscan: Hierarchical density based clustering. *Journal of Open Source Software* 2, 11 (2017), 205. <https://doi.org/10.21105/joss.00205>
- [12] Gestefane Rabbi, Júlia Viterbo, Gabriel Kakizaki, Zilton Cordeiro Junior, Raquel O. Prates, Julio C. S. Reis, and Marcos André Gonçalves. 2026. Detecting and Analysing Duplicate Consumer Complaints and Collective Demands Across Multiple Platforms. *Journal of Information and Data Management* (2026). <https://journals-sol.sbc.org.br/index.php/jidm/article/view/5959>
- [13] Daniyar Rakhimzhanov, Saule Belginova, and Didar Yedilkhan. 2025. Automated Classification of Public Transport Complaints via Text Mining Using LLMs and Embeddings. *Information* 16, 8 (2025), 644. <https://doi.org/10.3390/info16080644>
- [14] Nils Reimers and Iryna Gurevych. 2019. Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. In *Proceedings of EMNLP-IJCNLP 2019*. Association for Computational Linguistics, 3982–3992. <https://doi.org/10.18653/v1/D19-1410>

- [15] Julia Romberg and Tobias Escher. 2023. Making Sense of Citizens' Input through Artificial Intelligence: A Review of Methods for Computational Text Analysis to Support the Evaluation of Contributions in Public Participation. *Digital Government: Research and Practice* 4, 3, Article 26 (2023). <https://doi.org/10.1145/3603254>
- [16] Gerard Salton and Christopher Buckley. 1988. Term-Weighting Approaches in Automatic Text Retrieval. *Information Processing & Management* 24, 5 (1988), 513–523. [https://doi.org/10.1016/0306-4573\(88\)90021-0](https://doi.org/10.1016/0306-4573(88)90021-0)
- [17] Erich Schubert, Jörg Sander, Martin Ester, Hans-Peter Kriegel, and Xiaowei Xu. 2017. DBSCAN Revisited, Revisited: Why and How You Should (Still) Use DBSCAN. *ACM Transactions on Database Systems* 42, 3 (2017), 19:1–19:21. <https://doi.org/10.1145/3068335>
- [18] Ticiana L. Coelho da Silva, Celso França, Marcos André Gonçalves, Leonardo Rocha, Leonardo Alamy, Fernando Sola Pereira, and Eduardo Soares de Paiva. 2026. Bridging Citizens and Public Services: Improving Service Association with Retrieval-Augmented Generation (RAG) Labels. In *Proceedings of the 17th International Conference on Computational Processing of Portuguese (PROPOR 2026)*. Association for Computational Linguistics, 131–140. <https://aclanthology.org/2026.propor-1.13/>
- [19] Fábio Souza, Roberto Nogueira, and Roberto Lotufo. 2020. BERTimbau: Pretrained BERT Models for Brazilian Portuguese. In *Intelligent Systems (Lecture Notes in Computer Science, Vol. 12319)*. Springer, 403–417. https://doi.org/10.1007/978-3-030-61377-8_28
- [20] Zhi Tang, Xin Pan, and Zhiwei Gu. 2024. Analyzing Public Demands on China's Online Government Inquiry Platform: A BERTopic-Based Topic Modeling Study. *PLOS ONE* 19, 2 (2024), e0296855. <https://doi.org/10.1371/journal.pone.0296855>